

Fraud Detection Modeling for Credit Card Transactions:

A Machine Learning Approach

Francis Boateng

Abstract

This project investigates the effectiveness of machine learning models in identifying fraudulent transactions in a large dataset of over 550,000 records featuring European cardholder transactions from 2023. With scaled features and anonymized names for privacy, the dataset provides a robust foundation for model development. The goal is to train and evaluate various machine learning algorithms to distinguish legitimate from fraudulent transactions, assessing their accuracy and performance. The findings of this project will contribute to the development of efficient fraud detection systems, enhancing the security of financial transactions and minimizing potential losses for cardholders and financial institutions.

Exploratory Data Analysis

In this step, I will get to know the dataset better. I will look at summary statistics to understand the main trends and variations in the data. I'll also check the distribution of each variable to see if there are any patterns or unusual values. Additionally, I'll see how the variables are related, which will help me decide how to build my model. I'll examine how the variables behave differently depending on whether the transaction is normal or fraudulent. Finally, I'll check if there is an imbalance between normal and fraudulent transactions, which is essential for building accurate fraud detection models. By doing this, I will gain a deeper understanding of the data and identify significant trends and relationships to help me build a better model.

```
# Loading the dataset and removing the id column
fraud <- read.csv("creditcard_2023.csv")
FRAUD <- fraud[,-1]
head(FRAUD)
```

##	V1	V2	V3	V4	V5	V6	V7
----	----	----	----	----	----	----	----

```

## 1 -0.26064780 -0.4696485 2.4962661 -0.08372391 0.12968124 0.7328982 0.5190136
## 2 0.98509973 -0.3560451 0.5580564 -0.42965390 0.27714026 0.4286045 0.4064660
## 3 -0.26027161 -0.9493846 1.7285378 -0.45798629 0.07406165 1.4194811 0.7435111
## 4 -0.15215210 -0.5089587 1.7468401 -1.09017794 0.24948577 1.1433123 0.5182686
## 5 -0.20681952 -0.1652802 1.5270527 -0.44829266 0.10612511 0.5305489 0.6588491
## 6 0.02530229 -0.1405138 1.1911378 -0.70797881 0.43049032 0.4589732 0.6110496
##          V8          V9          V10          V11          V12          V13          V14
## 1 -0.13000605 0.7271593 0.6377345 -0.98702001 0.2934381 -0.9413861 0.5490199
## 2 -0.13311827 0.3474519 0.5298080 0.14010733 1.5642458 0.5740740 0.6277187
## 3 -0.09557601 -0.2612966 0.6907078 -0.27298493 0.6592007 0.8051732 0.6168744
## 4 -0.06512992 -0.2056976 0.5752307 -0.75258096 0.7374830 0.5929937 0.5595350
## 5 -0.21266001 1.0499208 0.9680461 -1.20317111 1.0295774 1.4393102 0.2414540
## 6 -0.09262861 0.1808114 0.4517884 0.03607131 0.8772389 -0.2897211 0.6309925
##          V15          V16          V17          V18          V19          V20          V21
## 1 1.8048786 0.21559799 0.5123067 0.3336437 0.1242702 0.0912019 -0.110551680
## 2 0.7061213 0.78918836 0.4038099 0.2017994 -0.3406871 -0.2339842 -0.194935964
## 3 3.0690248 -0.57751352 0.8865260 0.2394417 -2.3660789 0.3616523 -0.005020278
## 4 -0.6976637 -0.03066898 0.2426292 2.1786160 -1.3450602 -0.3782233 -0.146927137
## 5 0.1530079 0.22453813 0.3664662 0.2917816 0.4453167 0.2472370 -0.106984018
## 6 0.5602009 0.74113155 0.4217663 0.3625039 -0.2427488 -0.0764003 -0.187739355
##          V22          V23          V24          V25          V26          V27
## 1 0.21760614 -0.13479449 0.1659591 0.1262800 -0.4348240 -0.08123011
## 2 -0.60576091 0.07946908 -0.5773949 0.1900897 0.2965027 -0.24805206
## 3 0.70290638 0.94504549 -1.1546656 -0.6055637 -0.3128945 -0.30025804
## 4 -0.03821246 -0.21404819 -1.8931311 1.0039631 -0.5159503 -0.16531649
## 5 0.72972739 -0.16166570 0.3125610 -0.4141162 1.0711256 0.02371160
## 6 -0.53851811 -0.05046499 -0.6315531 -0.4564800 0.2526699 0.06668093
##          V28 Amount Class
## 1 -0.15104549 17982.10 0
## 2 -0.06451192 6531.37 0
## 3 -0.24471823 2513.54 0
## 4 0.04842363 5384.44 0
## 5 0.41911727 14278.97 0
## 6 0.09581151 6901.49 0

```

```

# Checking for missing values in the data set
colSums(is.na(FRAUD))

```

```

##      V1      V2      V3      V4      V5      V6      V7      V8      V9      V10     V11
##      0      0      0      0      0      0      0      0      0      0      0
##     V12     V13     V14     V15     V16     V17     V18     V19     V20     V21     V22
##      0      0      0      0      0      0      0      0      0      0      0
##     V23     V24     V25     V26     V27     V28 Amount  Class
##      0      0      0      0      0      0      0      0      0

```

```
# Summary Statistics
print(summary(FRAUD))
```

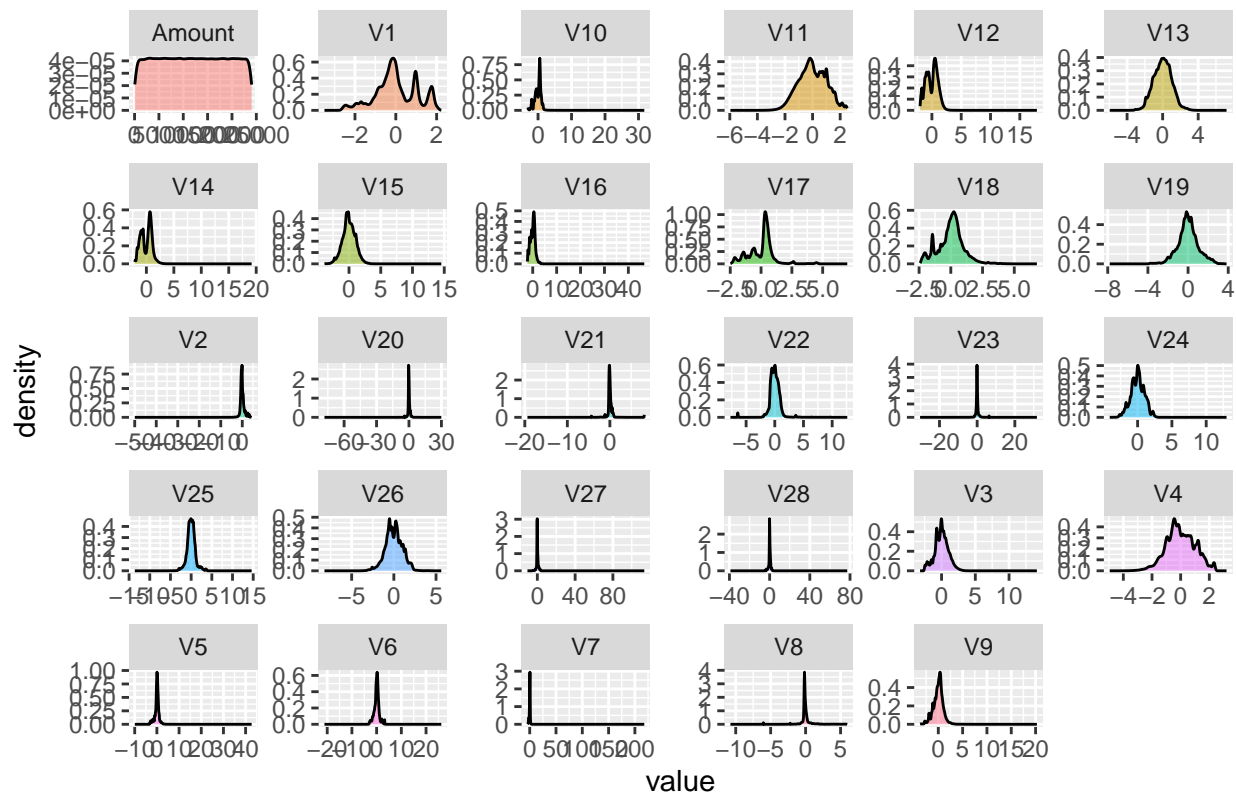
```
##          V1          V2          V3          V4
## Min.    :-3.49558  Min.    :-49.9666  Min.    :-3.183760  Min.    :-4.95122
## 1st Qu.: -0.56529  1st Qu.: -0.4867   1st Qu.: -0.649299  1st Qu.: -0.65602
## Median : -0.09364  Median : -0.1359   Median : 0.000353   Median : -0.07376
## Mean    : 0.00000   Mean    : 0.0000    Mean    : 0.000000   Mean    : 0.00000
## 3rd Qu.: 0.83266   3rd Qu.: 0.3436    3rd Qu.: 0.628538   3rd Qu.: 0.70700
## Max.    : 2.22905   Max.    : 4.3619    Max.    :14.125834   Max.    : 3.20154
##          V5          V6          V7          V8
## Min.    :-9.95279  Min.    :-21.11111  Min.    : -4.3518   Min.    :-10.7563
## 1st Qu.: -0.29350  1st Qu.: -0.44587  1st Qu.: -0.2835   1st Qu.: -0.1923
## Median : 0.08109   Median : 0.07872   Median : 0.2334    Median : -0.1145
## Mean    : 0.00000   Mean    : 0.00000   Mean    : 0.0000    Mean    : 0.0000
## 3rd Qu.: 0.43974   3rd Qu.: 0.49779   3rd Qu.: 0.5260    3rd Qu.: 0.0473
## Max.    :42.71689  Max.    : 26.16840  Max.    :217.8730   Max.    : 5.9580
##          V9          V10         V11         V12
## Min.    :-3.75192  Min.    :-3.1633   Min.    :-5.95472  Min.    :-2.0204
## 1st Qu.: -0.56874  1st Qu.: -0.5901   1st Qu.: -0.70145  1st Qu.: -0.8311
## Median : 0.09253   Median : 0.2626    Median : -0.04105  Median : 0.1621
## Mean    : 0.00000   Mean    : 0.0000    Mean    : 0.00000   Mean    : 0.0000
## 3rd Qu.: 0.55926   3rd Qu.: 0.5925    3rd Qu.: 0.74777   3rd Qu.: 0.7447
## Max.    :20.27006  Max.    :31.7227   Max.    : 2.51357   Max.    :17.9136
##          V13         V14         V15         V16
## Min.    :-5.95523  Min.    :-2.1074   Min.    :-3.86181  Min.    :-2.2145
## 1st Qu.: -0.69667  1st Qu.: -0.8732   1st Qu.: -0.62125  1st Qu.: -0.7163
## Median : 0.01761   Median : 0.2305    Median : -0.03926  Median : 0.1340
## Mean    : 0.00000   Mean    : 0.0000    Mean    : 0.00000   Mean    : 0.0000
## 3rd Qu.: 0.68561   3rd Qu.: 0.7518    3rd Qu.: 0.66541   3rd Qu.: 0.6556
## Max.    : 7.18749  Max.    :19.1695   Max.    :14.53220   Max.    :46.6529
##          V17         V18         V19         V20
## Min.    :-2.4849   Min.    :-2.42195  Min.    :-7.80499  Min.    :-78.1478
## 1st Qu.: -0.6195   1st Qu.: -0.55605  1st Qu.: -0.56531  1st Qu.: -0.3502
## Median : 0.2716    Median : 0.08729   Median : -0.02598  Median : -0.1234
## Mean    : 0.0000    Mean    : 0.00000   Mean    : 0.00000   Mean    : 0.0000
## 3rd Qu.: 0.5182    3rd Qu.: 0.54439   3rd Qu.: 0.56012   3rd Qu.: 0.2482
## Max.    : 6.9941    Max.    : 6.78372   Max.    : 3.83167   Max.    : 29.8728
##          V21         V22         V23         V24
## Min.    :-19.38252  Min.    :-7.73480  Min.    :-30.29545  Min.    :-4.0680
## 1st Qu.: -0.16644   1st Qu.: -0.49049  1st Qu.: -0.23763  1st Qu.: -0.6516
## Median : -0.03743   Median : -0.02733  Median : -0.05969  Median : 0.0159
## Mean    : 0.00000   Mean    : 0.00000   Mean    : 0.00000   Mean    : 0.0000
## 3rd Qu.: 0.14798    3rd Qu.: 0.46388   3rd Qu.: 0.15572   3rd Qu.: 0.7007
```

```
## Max.    : 8.08708    Max.    :12.63251    Max.    : 31.70763    Max.    :12.9656
##      V25              V26              V27              V28
## Min.    :-13.612633  Min.    :-8.22697   Min.    :-10.4986   Min.    :-39.03524
## 1st Qu.: -0.554148   1st Qu.: -0.63189   1st Qu.: -0.3050   1st Qu.: -0.23188
## Median : -0.008193   Median : -0.01189   Median : -0.1729   Median : -0.01393
## Mean    : 0.000000    Mean    : 0.00000    Mean    : 0.0000    Mean    : 0.00000
## 3rd Qu.: 0.550015    3rd Qu.: 0.67289    3rd Qu.: 0.3340    3rd Qu.: 0.40959
## Max.    : 14.621509   Max.    : 5.62329    Max.    :113.2311   Max.    : 77.25594
##      Amount          Class
## Min.    : 50.01      Min.    :0.0
## 1st Qu.: 6054.89     1st Qu.:0.0
## Median :12030.15     Median :0.5
## Mean    :12041.96     Mean    :0.5
## 3rd Qu.:18036.33     3rd Qu.:1.0
## Max.    :24039.93     Max.    :1.0
```

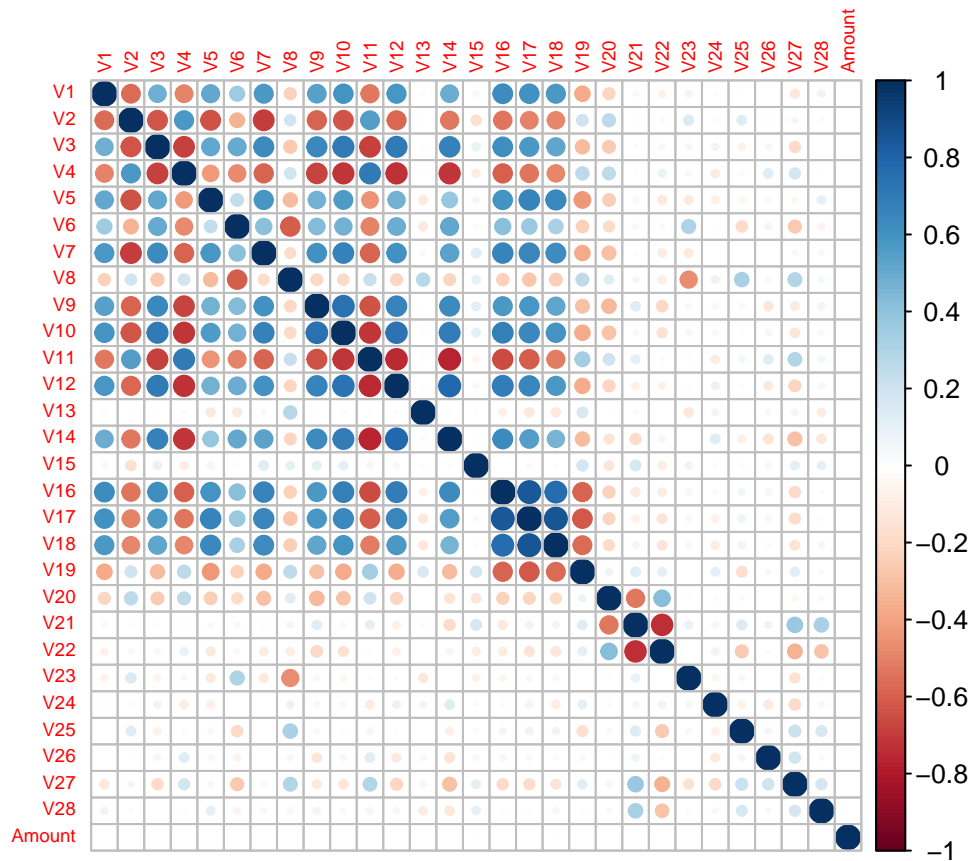
```
# Creating distribution plots for the variables
```

```
FRAUD %>%
  select(-Class) %>%
  gather() %>%
  ggplot(aes(value, fill = key)) +
  geom_density(alpha = 0.5) +
  facet_wrap(~key, scales = "free") +
  labs(title = "Distribution of Variables") +
  guides(fill = "none")
```

Distribution of Variables

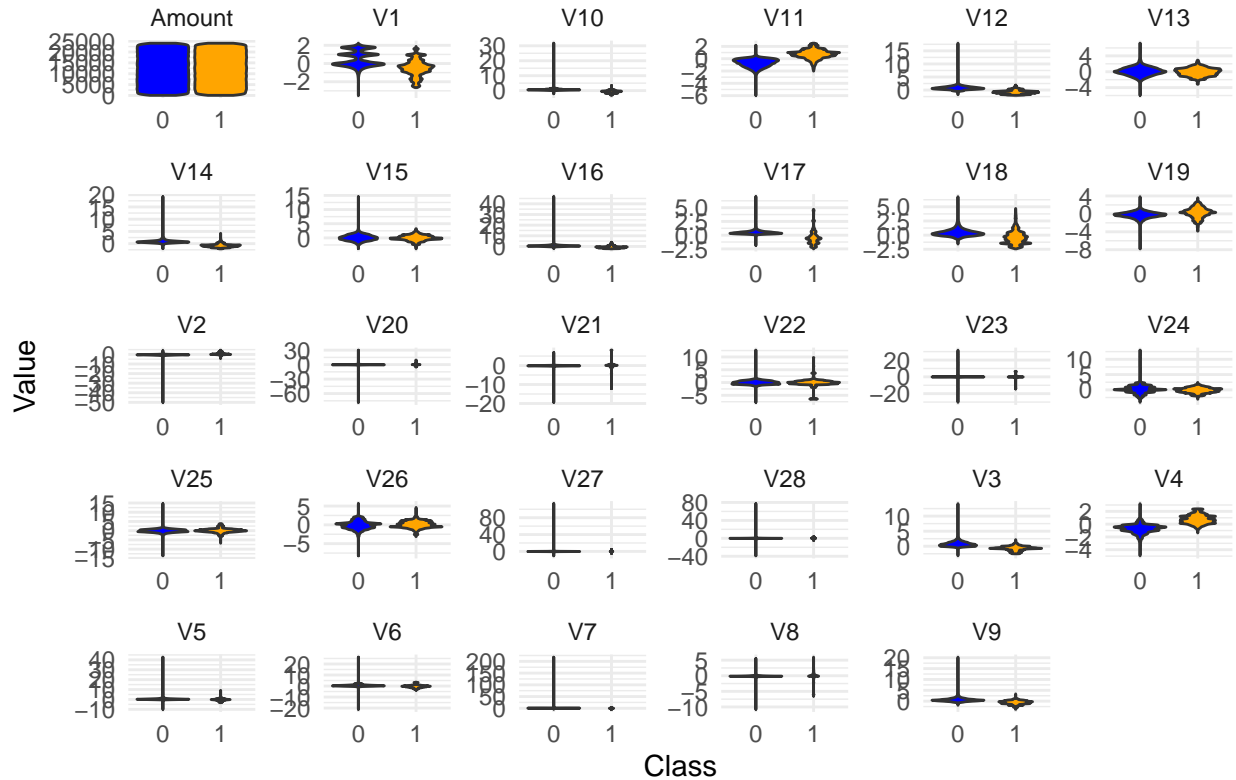


```
# Creating the correlation plot to check the correlation between the variables
corr_matrix <- FRAUD %>%
  select(-Class) %>%
  cor()
corrplot(corr_matrix, method = "circle", tl.cex = 0.6)
```

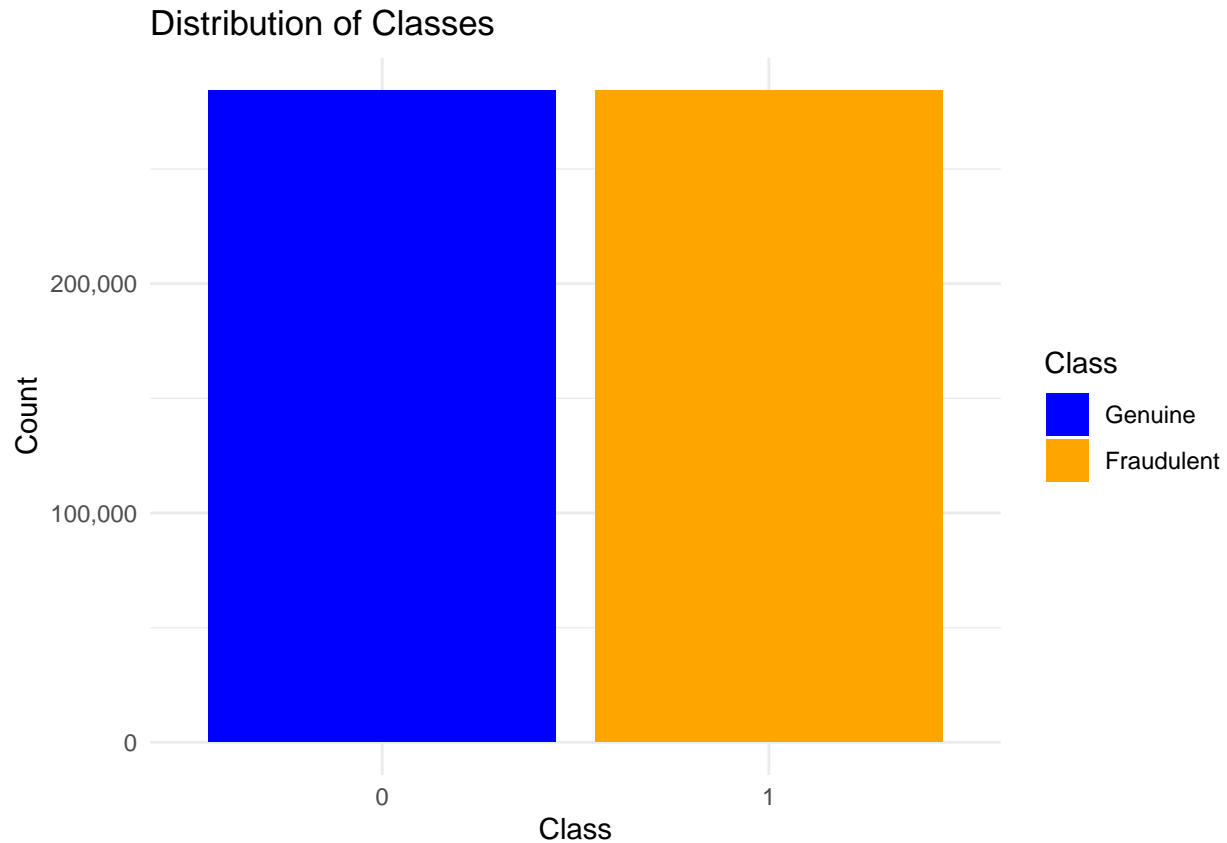


```
# Creating a violin plot distribution by class
FRAUD %>% gather(key = "Variable", value = "Value", -Class) %>%
  mutate(Class = as.factor(Class)) %>%
  ggplot(aes(x = Class, y = Value, fill = Class)) +
  geom_violin() +
  facet_wrap(~ Variable, scales = "free") +
  labs(title = "Distribution of Variables by Class",
       x = "Class",
       y = "Value") +
  scale_fill_manual(values = c("0" = "blue", "1" = "orange")) +
  theme_minimal() +
  guides(fill = "none")
```

Distribution of Variables by Class



```
# Plotting the class distribution
ggplot(FRAUD, aes(x = factor(Class), fill = factor(Class))) +
  geom_bar() +
  labs(x = "Class", y = "Count", fill = "Class", title = "Distribution of Classes") +
  scale_fill_manual(values = c("0" = "blue", "1" = "orange"), labels = c("Genuine", "Fraud")) +
  scale_y_continuous(labels = scales::comma) +
  theme_minimal()
```



By exploring the data, I found some interesting things. Some variables are very closely related, meaning they tend to change together in a predictable way, while others move in opposite directions. When I looked at the patterns of fraudulent and genuine transactions, I saw some clear differences that could help me tell them apart.

I also noticed that the number of fraudulent and genuine transactions is almost equal, which is unusual because in real life, there are usually many more genuine transactions than fraudulent ones. This suggests that the data was created artificially, rather than coming from real-world situations. As a result, I won't need to use special techniques to balance the data before training my model.

Data Prepaion for Model Building

I'll now prepare the dataset for model training and testing. I'll normalize the 'Amount' feature to be on the same scale as the other variables. This is because it's the only feature that still needs to be transformed using PCA.

Then, I'll divide the dataset into two parts: training and testing sets. I'll use 70% of the data to train the models and the remaining 30% to test their performance. This way, I can ensure the models are well-trained and test them thoroughly in a controlled environment.


```
# Normalizing the 'Amount' variable
preProcessRange <- preProcess(FRAUD["Amount"], method = c("center", "scale"))
FRAUD_norm <- predict(preProcessRange, FRAUD["Amount"])
FRAUD <- bind_cols(FRAUD[, -which(names(FRAUD) %in% "Amount")], FRAUD_norm)
head(FRAUD)
```

```
##           V1           V2           V3           V4           V5           V6           V7
## 1 -0.26064780 -0.4696485  2.4962661 -0.08372391  0.12968124  0.7328982  0.5190136
## 2  0.98509973 -0.3560451  0.5580564 -0.42965390  0.27714026  0.4286045  0.4064660
## 3 -0.26027161 -0.9493846  1.7285378 -0.45798629  0.07406165  1.4194811  0.7435111
## 4 -0.15215210 -0.5089587  1.7468401 -1.09017794  0.24948577  1.1433123  0.5182686
## 5 -0.20681952 -0.1652802  1.5270527 -0.44829266  0.10612511  0.5305489  0.6588491
## 6  0.02530229 -0.1405138  1.1911378 -0.70797881  0.43049032  0.4589732  0.6110496
##           V8           V9           V10          V11          V12          V13          V14
## 1 -0.13000605  0.7271593  0.6377345 -0.98702001  0.2934381 -0.9413861  0.5490199
## 2 -0.13311827  0.3474519  0.5298080  0.14010733  1.5642458  0.5740740  0.6277187
## 3 -0.09557601 -0.2612966  0.6907078 -0.27298493  0.6592007  0.8051732  0.6168744
## 4 -0.06512992 -0.2056976  0.5752307 -0.75258096  0.7374830  0.5929937  0.5595350
## 5 -0.21266001  1.0499208  0.9680461 -1.20317111  1.0295774  1.4393102  0.2414540
## 6 -0.09262861  0.1808114  0.4517884  0.03607131  0.8772389 -0.2897211  0.6309925
##           V15          V16          V17          V18          V19          V20          V21
## 1  1.8048786  0.21559799  0.5123067  0.3336437  0.1242702  0.0912019 -0.110551680
## 2  0.7061213  0.78918836  0.4038099  0.2017994 -0.3406871 -0.2339842 -0.194935964
## 3  3.0690248 -0.57751352  0.8865260  0.2394417 -2.3660789  0.3616523 -0.005020278
## 4 -0.6976637 -0.03066898  0.2426292  2.1786160 -1.3450602 -0.3782233 -0.146927137
## 5  0.1530079  0.22453813  0.3664662  0.2917816  0.4453167  0.2472370 -0.106984018
## 6  0.5602009  0.74113155  0.4217663  0.3625039 -0.2427488 -0.0764003 -0.187739355
##           V22          V23          V24          V25          V26          V27
## 1  0.21760614 -0.13479449  0.1659591  0.1262800 -0.4348240 -0.08123011
## 2 -0.60576091  0.07946908 -0.5773949  0.1900897  0.2965027 -0.24805206
## 3  0.70290638  0.94504549 -1.1546656 -0.6055637 -0.3128945 -0.30025804
## 4 -0.03821246 -0.21404819 -1.8931311  1.0039631 -0.5159503 -0.16531649
## 5  0.72972739 -0.16166570  0.3125610 -0.4141162  1.0711256  0.02371160
## 6 -0.53851811 -0.05046499 -0.6315531 -0.4564800  0.2526699  0.06668093
##           V28 Class      Amount
## 1 -0.15104549      0  0.8584462
## 2 -0.06451192      0 -0.7963686
## 3 -0.24471823      0 -1.3770097
## 4  0.04842363      0 -0.9621185
## 5  0.41911727      0  0.3232843
## 6  0.09581151      0 -0.7428803
```

```
# Splitting the data into training and test sets (70% training, 30% test)
set.seed(123)
split <- sample.split(FRAUD$Class, SplitRatio = 0.7)
training_set <- subset(FRAUD, split == TRUE)
test_set <- subset(FRAUD, split == FALSE)

nrow(training_set)
```

```
## [1] 398040
```

```
nrow(test_set)
```

```
## [1] 170590
```

Developing a Statistical Model for Predictive Analytics

I'll now create four different machine-learning models to help identify fraudulent transactions. These models are logistic regression, decision trees, random forests, and a simple neural network.

Next, I'll test how well these models work by using two important measures: the Area Under the Precision-Recall Curve (AUPRC) and the Area Under the ROC Curve (AUROC). This will help me see how well each model detects fraudulent transactions.

1. Logistic Regression

```
set.seed(123)
glm_model = glm(Class ~ ., data = training_set, family = binomial(link = 'logit'))
summary(glm_model)
```

```
##
## Call:
## glm(formula = Class ~ ., family = binomial(link = "logit"), data = training_set)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -8.490  -0.141   0.000   0.000   3.889
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  9.043530  0.098505  91.808 < 2e-16 ***
```

```
## V1          -0.686264  0.022348 -30.708 < 2e-16 ***
## V2           0.176351  0.019680   8.961 < 2e-16 ***
## V3          -1.182274  0.023568 -50.163 < 2e-16 ***
## V4           3.634341  0.032042 113.424 < 2e-16 ***
## V5          -0.002097  0.017339  -0.121  0.904
## V6          -0.476370  0.019753 -24.117 < 2e-16 ***
## V7          -1.089877  0.031219 -34.910 < 2e-16 ***
## V8          -2.805206  0.053037 -52.892 < 2e-16 ***
## V9          -0.465499  0.026723 -17.419 < 2e-16 ***
## V10         -1.899152  0.037673 -50.412 < 2e-16 ***
## V11          1.872530  0.021693  86.319 < 2e-16 ***
## V12         -2.798026  0.030775 -90.919 < 2e-16 ***
## V13          0.009754  0.010479   0.931  0.352
## V14         -3.355914  0.031563 -106.324 < 2e-16 ***
## V15         -0.241435  0.010313 -23.411 < 2e-16 ***
## V16         -0.807547  0.029040 -27.808 < 2e-16 ***
## V17         -1.935311  0.032241 -60.026 < 2e-16 ***
## V18         -0.949279  0.024826 -38.238 < 2e-16 ***
## V19         -0.071113  0.014976  -4.748 2.05e-06 ***
## V20          0.150098  0.014699  10.212 < 2e-16 ***
## V21          0.255312  0.033909   7.529 5.10e-14 ***
## V22          0.447060  0.018055  24.761 < 2e-16 ***
## V23         -0.331602  0.012987 -25.533 < 2e-16 ***
## V24         -0.159295  0.010990 -14.494 < 2e-16 ***
## V25          0.173847  0.012963  13.411 < 2e-16 ***
## V26         -0.110695  0.011690  -9.469 < 2e-16 ***
## V27          0.194615  0.026363   7.382 1.56e-13 ***
## V28          0.147788  0.012256  12.059 < 2e-16 ***
## Amount      0.007752  0.009629   0.805  0.421
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 551801  on 398039  degrees of freedom
## Residual deviance:  74941  on 398010  degrees of freedom
## AIC: 75001
##
## Number of Fisher Scoring iterations: 25
```

```
pred_glm <- predict(glm_model, newdata = test_set, type = 'response')
```

```
# ROC and PR Curves for GLM model
glm_fg <- pred_glm[test_set$Class == 1]
```


```

glm_bg <- pred_glm[test_set$Class == 0]

glm_roc <- roc.curve(scores.class0 = glm_fg , scores.class1 = glm_bg, curve = T)
glm_pr <- pr.curve(scores.class0 = glm_fg , scores.class1 = glm_bg, curve = T)

options(repr.plot.width=12, repr.plot.height=7)
par(mfrow = c(1, 2))
plot(glm_roc, col = "blue", main = "ROC Curve")
plot(glm_pr, col = "blue", main = "PR Curve")

```



The Logistic Regression model showed AUROC of 0.9935217 and AUPRC of 0.9946013.

2. Decision Tree

```

set.seed(123)
dt_model <- rpart(Class ~ ., data = training_set, method = "class")
summary(dt_model)

```

Call:

```

## rpart(formula = Class ~ ., data = training_set, method = "class")
##   n= 398040
##
##           CP nsplit rel error   xerror   xstd
## 1 0.8637825      0 1.0000000 1.0039092 0.0015850148
## 2 0.0100000      1 0.1362175 0.1365642 0.0007995806
##
## Variable importance
## V14 V10 V12 V11 V17  V3
##  20  17  16  16  16  15
##
## Node number 1: 398040 observations,   complexity param=0.8637825
##   predicted class=0   expected loss=0.5   P(node) =1
##   class counts: 199020 199020
##   probabilities: 0.500 0.500
##   left son=2 (216200 obs) right son=3 (181840 obs)
##   Primary splits:
##     V14 < 0.008223371 to the right, improve=149607.7, (0 missing)
##     V10 < 0.0121387   to the right, improve=138591.8, (0 missing)
##     V12 < -0.2313897  to the right, improve=126566.6, (0 missing)
##     V4  < -0.03006013 to the left,  improve=122499.8, (0 missing)
##     V17 < 0.09257648  to the right, improve=122089.1, (0 missing)
##   Surrogate splits:
##     V10 < 0.01730558  to the right, agree=0.948, adj=0.886, (0 split)
##     V12 < -0.2313897  to the right, agree=0.918, adj=0.821, (0 split)
##     V11 < 0.2068063   to the left,  agree=0.918, adj=0.820, (0 split)
##     V17 < 0.09264274  to the right, agree=0.913, adj=0.809, (0 split)
##     V3  < -0.1272986  to the right, agree=0.883, adj=0.744, (0 split)
##
## Node number 2: 216200 observations
##   predicted class=0   expected loss=0.1024283   P(node) =0.5431615
##   class counts: 194055 22145
##   probabilities: 0.898 0.102
##
## Node number 3: 181840 observations
##   predicted class=1   expected loss=0.02730422   P(node) =0.4568385
##   class counts:  4965 176875
##   probabilities: 0.027 0.973

```

```

pred_dt <- predict(dt_model, newdata = test_set, type = "prob")

```

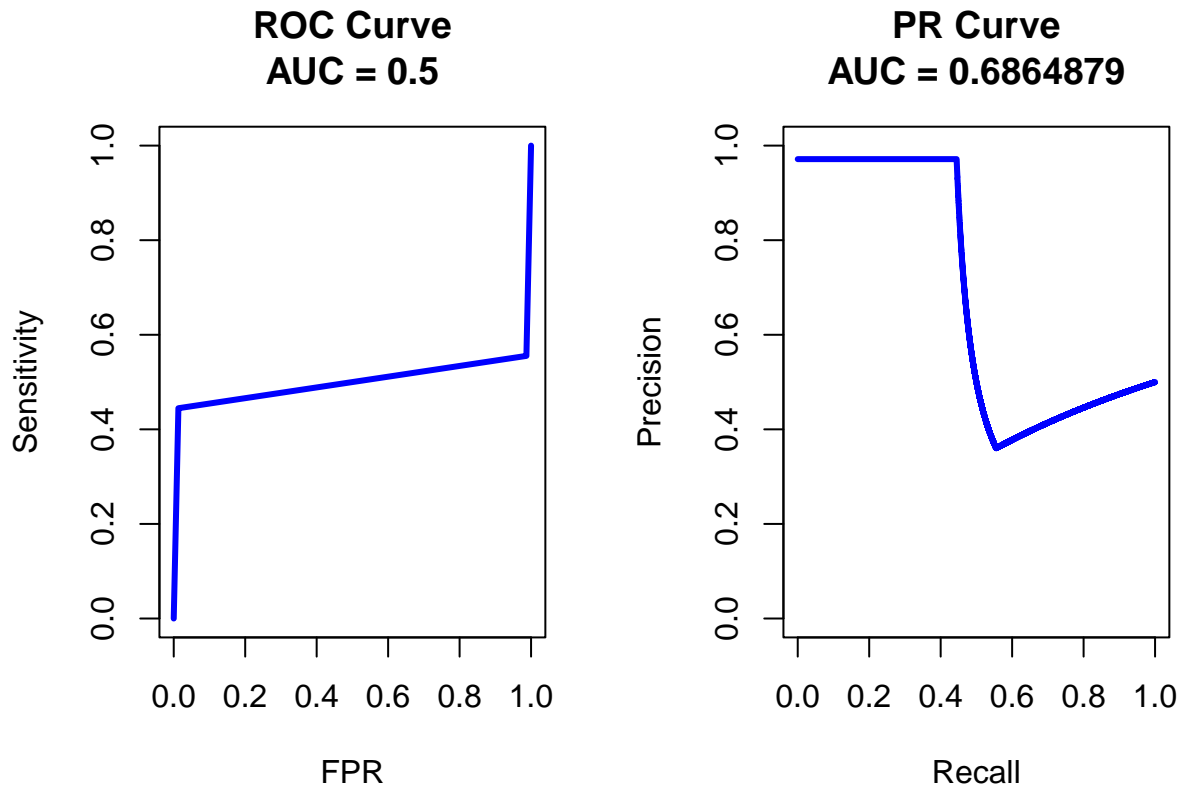
```

# ROC and PR Curves for decision tree model
dt_fg <- pred_dt[test_set$Class == 1]
dt_bg <- pred_dt[test_set$Class == 0]

```

```
dt_roc <- roc.curve(scores.class0 = dt_fg , scores.class1 = dt_bg, curve = T)
dt_pr <- pr.curve(scores.class0 = dt_fg , scores.class1 = dt_bg, curve = T)

par(mfrow = c(1, 2))
plot(dt_roc, col = "blue", main = "ROC Curve")
plot(dt_pr, col = "blue", main = "PR Curve")
```



The Decision Tree model showed AUROC of 0.5 and AUPRC of 0.6864879.

3. Random Forest

```
set.seed(123)
rf_model <- randomForest(Class ~ ., data = training_set, ntree = 10)
```

```
## Warning in randomForest.default(m, y, ...): The response has five or fewer
## unique values. Are you sure you want to do regression?
```

```
summary(rf_model)
```

```
##              Length Class  Mode
## call              4 -none- call
## type              1 -none- character
## predicted        398040 -none- numeric
## mse               10 -none- numeric
## rsq               10 -none- numeric
## oob.times        398040 -none- numeric
## importance        29 -none- numeric
## importanceSD       0 -none- NULL
## localImportance    0 -none- NULL
## proximity         0 -none- NULL
## ntree             1 -none- numeric
## mtry              1 -none- numeric
## forest            11 -none- list
## coefs             0 -none- NULL
## y                398040 -none- numeric
## test              0 -none- NULL
## inbag             0 -none- NULL
## terms             3 terms  call
```

```
pred_rf <- predict(rf_model, newdata = test_set, type = "class")
```

```
# ROC and PR Curves for random forest model
```

```
rf_fg <- pred_rf[test_set$Class == 1]
```

```
rf_bg <- pred_rf[test_set$Class == 0]
```


```
rf_roc <- roc.curve(scores.class0 = rf_fg , scores.class1 = rf_bg, curve = T)
```

```
rf_pr <- pr.curve(scores.class0 = rf_fg , scores.class1 = rf_bg, curve = T)
```

```
par(mfrow = c(1, 2))
```

```
plot(rf_roc, col = "blue", main = "ROC Curve")
```

```
plot(rf_pr, col = "blue", main = "PR Curve")
```



The Random Forest model showed AUROC of 0.9999996 and AUPRC of 0.9999995.

4. Neural Network

```
set.seed(123)
nn_model <- nnet(Class ~ ., data = training_set, size = 10, linout = FALSE, maxit = 200)
```

```
## # weights: 311
## initial value 102009.589037
## iter 10 value 34177.761401
## iter 20 value 24791.985455
## iter 30 value 20678.775000
## iter 40 value 20465.142281
## iter 50 value 19403.995463
## iter 60 value 19247.385257
## iter 70 value 18598.562075
## iter 80 value 18299.518540
## iter 90 value 18257.028098
## iter 100 value 18118.303157
## iter 110 value 17824.777563
```



```

## iter 120 value 17715.848363
## iter 130 value 17526.599233
## iter 140 value 17381.762525
## iter 150 value 17371.795110
## iter 160 value 17354.815553
## iter 170 value 17299.419974
## iter 180 value 17227.509980
## iter 190 value 17171.664654
## iter 200 value 17144.778410
## final value 17144.778410
## stopped after 200 iterations

```

```
summary(nn_model)
```

```

## a 29-10-1 network with 311 weights
## options were -
##  b->h1  i1->h1  i2->h1  i3->h1  i4->h1  i5->h1  i6->h1  i7->h1  i8->h1  i9->h1
##  36.07   69.87  -56.45   84.33  -90.61   45.56   55.64   57.26  -24.29   66.09
## i10->h1 i11->h1 i12->h1 i13->h1 i14->h1 i15->h1 i16->h1 i17->h1 i18->h1 i19->h1
##  80.43  -87.45   93.90   -4.06   98.10  -16.21   73.68   63.76   56.49  -42.18
## i20->h1 i21->h1 i22->h1 i23->h1 i24->h1 i25->h1 i26->h1 i27->h1 i28->h1 i29->h1
## -22.11  -17.28   -5.83   -2.99    4.25   -8.79   -3.59  -22.05   -6.19    9.30
##  b->h2  i1->h2  i2->h2  i3->h2  i4->h2  i5->h2  i6->h2  i7->h2  i8->h2  i9->h2
##  54.15   49.25  131.86 -129.91  217.28 -192.38  102.10   13.00  -41.57 -121.43
## i10->h2 i11->h2 i12->h2 i13->h2 i14->h2 i15->h2 i16->h2 i17->h2 i18->h2 i19->h2
## -122.96  120.72 -275.77  -60.87 -268.02 -100.08  102.82  117.60  -69.52  -98.05
## i20->h2 i21->h2 i22->h2 i23->h2 i24->h2 i25->h2 i26->h2 i27->h2 i28->h2 i29->h2
## -22.14   55.41 -162.66 -121.61 -112.03  228.38   89.89   29.21   29.46   17.02
##  b->h3  i1->h3  i2->h3  i3->h3  i4->h3  i5->h3  i6->h3  i7->h3  i8->h3  i9->h3
## 343.41   93.61   47.99 -259.34  112.10  -37.53   47.87   78.08    6.54   44.55
## i10->h3 i11->h3 i12->h3 i13->h3 i14->h3 i15->h3 i16->h3 i17->h3 i18->h3 i19->h3
## -191.26  191.94 -110.57   79.27 -349.74 -237.74  -76.20  130.42  137.13 -117.00
## i20->h3 i21->h3 i22->h3 i23->h3 i24->h3 i25->h3 i26->h3 i27->h3 i28->h3 i29->h3
## -129.78   28.89  158.57   71.94   62.95 -194.33 -405.25   -2.75  101.09  271.12
##  b->h4  i1->h4  i2->h4  i3->h4  i4->h4  i5->h4  i6->h4  i7->h4  i8->h4  i9->h4
## -210.89   72.20  -59.82  129.27 -128.92   33.14   98.24   58.78  -29.31   99.31
## i10->h4 i11->h4 i12->h4 i13->h4 i14->h4 i15->h4 i16->h4 i17->h4 i18->h4 i19->h4
##  129.62 -155.84  165.76    0.52  186.84  -42.16  135.08  112.20   59.59  -52.70
## i20->h4 i21->h4 i22->h4 i23->h4 i24->h4 i25->h4 i26->h4 i27->h4 i28->h4 i29->h4
## -53.15  -34.08   10.55   11.97    9.95   24.50   -8.41  -73.49  -53.46   11.23
##  b->h5  i1->h5  i2->h5  i3->h5  i4->h5  i5->h5  i6->h5  i7->h5  i8->h5  i9->h5
## -170.52   64.14  -22.78  119.64 -116.99   15.27   56.00   42.74   39.80   55.83
## i10->h5 i11->h5 i12->h5 i13->h5 i14->h5 i15->h5 i16->h5 i17->h5 i18->h5 i19->h5
##  113.52 -164.15  172.56   22.97  183.53   34.80  131.32  138.12   85.96  -14.61

```

```

## i20->h5 i21->h5 i22->h5 i23->h5 i24->h5 i25->h5 i26->h5 i27->h5 i28->h5 i29->h5
## -14.58 -58.02 42.76 -36.38 -2.81 46.34 40.22 -44.00 -33.35 20.11
## b->h6 i1->h6 i2->h6 i3->h6 i4->h6 i5->h6 i6->h6 i7->h6 i8->h6 i9->h6
## -135.26 60.28 43.30 18.30 -158.50 53.18 -38.42 61.25 85.88 76.33
## i10->h6 i11->h6 i12->h6 i13->h6 i14->h6 i15->h6 i16->h6 i17->h6 i18->h6 i19->h6
## 61.29 -139.39 109.10 125.98 208.59 -121.09 37.66 66.14 25.22 -4.70
## i20->h6 i21->h6 i22->h6 i23->h6 i24->h6 i25->h6 i26->h6 i27->h6 i28->h6 i29->h6
## -22.93 10.71 -112.80 -21.21 -15.28 -93.50 -19.75 31.13 -5.26 10.83
## b->h7 i1->h7 i2->h7 i3->h7 i4->h7 i5->h7 i6->h7 i7->h7 i8->h7 i9->h7
## -95.33 162.23 22.28 116.16 -111.47 136.74 72.73 -65.38 13.25 29.74
## i10->h7 i11->h7 i12->h7 i13->h7 i14->h7 i15->h7 i16->h7 i17->h7 i18->h7 i19->h7
## 165.07 -2.28 169.07 -30.17 144.57 24.97 248.62 208.23 66.01 -80.06
## i20->h7 i21->h7 i22->h7 i23->h7 i24->h7 i25->h7 i26->h7 i27->h7 i28->h7 i29->h7
## 76.42 -50.07 -119.57 8.50 -8.15 2.04 37.70 -39.54 -41.87 174.51
## b->h8 i1->h8 i2->h8 i3->h8 i4->h8 i5->h8 i6->h8 i7->h8 i8->h8 i9->h8
## 124.65 -121.36 -66.42 -91.95 -38.00 32.87 -28.32 67.47 -32.37 5.10
## i10->h8 i11->h8 i12->h8 i13->h8 i14->h8 i15->h8 i16->h8 i17->h8 i18->h8 i19->h8
## -97.68 39.67 -86.34 9.90 -115.50 37.20 48.99 50.09 55.77 -84.45
## i20->h8 i21->h8 i22->h8 i23->h8 i24->h8 i25->h8 i26->h8 i27->h8 i28->h8 i29->h8
## -32.95 36.92 -4.23 -36.76 -104.81 85.92 109.39 0.34 -12.55 -0.72
## b->h9 i1->h9 i2->h9 i3->h9 i4->h9 i5->h9 i6->h9 i7->h9 i8->h9 i9->h9
## -298.18 71.28 52.50 169.27 -163.22 40.04 32.21 -37.27 40.17 42.86
## i10->h9 i11->h9 i12->h9 i13->h9 i14->h9 i15->h9 i16->h9 i17->h9 i18->h9 i19->h9
## 144.94 -99.53 189.55 25.80 178.69 -16.65 135.96 98.99 32.11 -74.66
## i20->h9 i21->h9 i22->h9 i23->h9 i24->h9 i25->h9 i26->h9 i27->h9 i28->h9 i29->h9
## 29.55 -73.92 64.25 -52.59 -40.95 12.49 80.74 -119.47 -13.32 14.37
## b->h10 i1->h10 i2->h10 i3->h10 i4->h10 i5->h10 i6->h10 i7->h10
## 167.83 -58.48 11.27 -121.24 156.36 -18.19 -7.46 -2.06
## i8->h10 i9->h10 i10->h10 i11->h10 i12->h10 i13->h10 i14->h10 i15->h10
## -37.84 -76.36 36.29 79.75 -156.68 93.96 -92.71 15.08
## i16->h10 i17->h10 i18->h10 i19->h10 i20->h10 i21->h10 i22->h10 i23->h10
## -103.94 -139.67 -114.53 85.75 5.06 -11.53 105.26 38.60
## i24->h10 i25->h10 i26->h10 i27->h10 i28->h10 i29->h10
## 47.39 -102.69 68.28 -6.92 -32.39 26.50
## b->o h1->o h2->o h3->o h4->o h5->o h6->o h7->o h8->o h9->o
## 335.70 -165.74 571.94 616.76 61.80 -160.02 -363.83 -455.47 292.62 -194.20
## h10->o
## 323.28

```

```

nn_predictions_probs <- predict(nn_model, newdata=test_set[, !names(test_set) %in% "Class"])
pred_nn <- ifelse(nn_predictions_probs > 0.5, 1, 0)

```

```

# ROC and PR Curves for random forest model
nn_fg <- pred_nn[test_set$Class == 1]

```

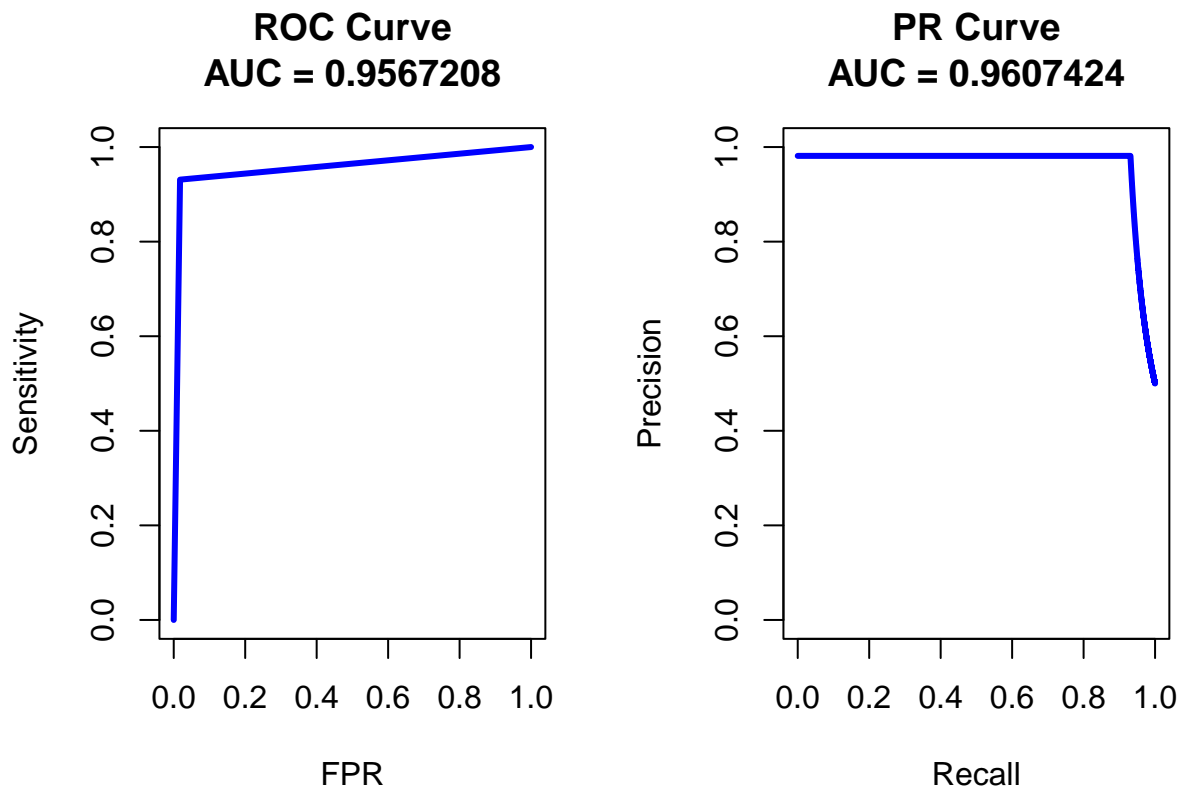
```

nn_bg <- pred_nn[test_set$Class == 0]

nn_roc <- roc.curve(scores.class0 = nn_fg , scores.class1 = nn_bg, curve = T)
nn_pr <- pr.curve(scores.class0 = nn_fg , scores.class1 = nn_bg, curve = T)

par(mfrow = c(1, 2))
plot(nn_roc, col = "blue", main = "ROC Curve")
plot(nn_pr, col = "blue", main = "PR Curve")

```



The Neural Network model showed AUROC of 0.9491647 and AUPRC of 0.9487003.

Model Evaluation and Comparison

Next, I'll check how well the four machine-learning models work. I'll use plots to show the results for each model, looking at how good they are at finding fraudulent transactions (AUPRC) and how well they can tell apart real and fake transactions (AUROC). These plots will help me see which model is best at finding fraud and how good it is at not mistaking genuine transactions for fake ones. By looking at these numbers, I'll find the model that's best at detecting fraud.

```

par(mfrow = c(1, 2))

# Extract ROC data for all models
roc_data_glm <- data.frame(FPR = glm_roc$curve[, 1], TPR = glm_roc$curve[, 2])
roc_data_dt <- data.frame(FPR = dt_roc$curve[, 1], TPR = dt_roc$curve[, 2])
roc_data_rf <- data.frame(FPR = rf_roc$curve[, 1], TPR = rf_roc$curve[, 2])
roc_data_nn <- data.frame(FPR = nn_roc$curve[, 1], TPR = nn_roc$curve[, 2])


# Plot ROC curves
plot(roc_data_glm$FPR, roc_data_glm$TPR, type = 'l', col = "orange", lwd = 5, xlab = "Fa
lines(roc_data_dt$FPR, roc_data_dt$TPR, col = "blue", lwd = 5)
lines(roc_data_rf$FPR, roc_data_rf$TPR, col = "brown", lwd = 5)
lines(roc_data_nn$FPR, roc_data_nn$TPR, col = "purple", lwd = 5)
legend("bottomright", legend = c("GLM", "Decision Tree", "Random Forest", "Neural Network

# Extract PR data for all models
pr_data_glm <- data.frame(Recall = glm_pr$curve[, 1], Precision = glm_pr$curve[, 2])
pr_data_dt <- data.frame(Recall = dt_pr$curve[, 1], Precision = dt_pr$curve[, 2])
pr_data_rf <- data.frame(Recall = rf_pr$curve[, 1], Precision = rf_pr$curve[, 2])
pr_data_nn <- data.frame(Recall = nn_pr$curve[, 1], Precision = nn_pr$curve[, 2])

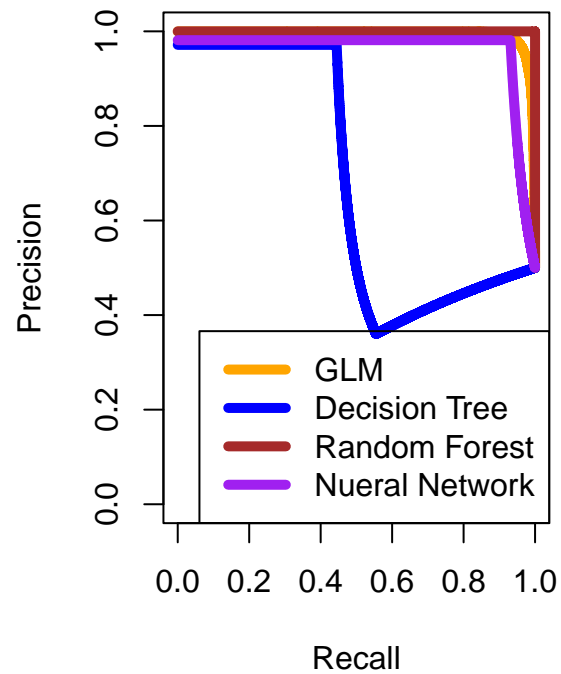
# Plot PR curves
plot(pr_data_glm$Recall, pr_data_glm$Precision, type = 'l', col = "orange", , lwd = 5, x
lines(pr_data_dt$Recall, pr_data_dt$Precision, col = "blue", lwd = 5)
lines(pr_data_rf$Recall, pr_data_rf$Precision, col = "brown", lwd = 5)
lines(pr_data_nn$Recall, pr_data_nn$Precision, col = "purple", lwd = 5)
legend("bottomright", legend = c("GLM", "Decision Tree", "Random Forest", "Nueral Network

```

ROC Curves Comparison



PR Curves Comparison



```
# Reset the plotting parameters to default
par(mfrow = c(1, 1))
```

```
# Extre AUC ROC values
```

```
auc_roc_glm <- glm_roc$auc
auc_roc_dt <- dt_roc$auc
auc_roc_rf <- rf_roc$auc
auc_roc_nn <- nn_roc$auc
```

```
# Extre AUC PR values
```

```
auc_pr_glm <- glm_pr$auc.integral
auc_pr_dt <- dt_pr$auc.integral
auc_pr_rf <- rf_pr$auc.integral
auc_pr_nn <- nn_pr$auc.integral
```

```
# Create a data frame
```

```
auc_table <- data.frame(
  Model = c("GLM", "Decision Tree", "Random Forest", "Neural Network"),
  AUC_ROC = c(auc_roc_glm, auc_roc_dt, auc_roc_rf, auc_roc_nn),
  AUC_PR = c(auc_pr_glm, auc_pr_dt, auc_pr_rf, auc_pr_nn)
)
```

```
# Print the table
```

```
print(auc_table)
```

```
##           Model    AUC_ROC    AUC_PR
## 1           GLM 0.9935217 0.9946013
## 2 Decision Tree 0.5000000 0.6864879
## 3 Random Forest 0.9999697 0.9999397
## 4 Neural Network 0.9567208 0.9607424
```

```
# Create the bar plot for the results
```


```
par(mfrow = c(1, 2))
```

```
# Create bar plot for AUROC
```

```
barplot(auc_table$AUC_ROC, names.arg = auc_table$Model, col = c("blue", "orange", "brown", "purple"))
```

```
# Create bar plot for AUPRC
```

```
barplot(auc_table$AUC_PR, names.arg = auc_table$Model, col = c("blue", "orange", "brown", "purple"))
```



```
# Reset graphical parameters to default  
par(mfrow = c(1, 1))
```

Conclusion

We wanted to find the best models for detecting credit card fraud using a dataset with fraudulent and genuine transactions. I developed four machine learning models: Logistic Regression, Decision Trees, Random Forest, and Neural Network.

The Random Forest model was almost perfect at telling apart genuine and fraudulent transactions, scoring 0.9999996. Logistic Regression was also perfect, with a score of 0.9935217, followed closely by Neural Network at 0.9491647. Decision Tree could have done better, with a score of 0.5, which means it struggled to differentiate between real and fake transactions.

When we looked at how well the models handled imbalanced data, Random Forest again did exceptionally well, scoring 0.9999995, showing it's very effective at detecting fraud. Logistic Regression was also robust, with a score of 0.9946013. Neural Network was decent, scoring 0.9487003, while Decision Tree struggled again, with a score of 0.6864879.

Overall, Random Forest was the most accurate model.